

Interview Arena

Briefing for the next AI agent

Consolidates the full planning conversation with the owner and a read-only review of the GitHub repo [rajeshkrogmax/YZIWORKSADMIN](#). Written 28 September 2026. Where this document and the current repo disagree, the repo wins: verify before relying on anything here.

0. Ground rules — read before anything else

1. Your job right now is to understand the project, not to build it. Write no code, commit nothing, push nothing and open no pull requests until the owner says so in chat.
2. The live site is production. The existing Meet AI Sera (the open trial) must keep working exactly as it does today. Do not edit its files.
3. The repo's HANDOFF.md says the dashboard is owned by another person and must not be touched. Follow that. Section 14 lists an open question about which codebase "the dashboard" refers to.
4. Read the repo in this order: docs/HANDOFF.md, CLAUDE.md, netlify/functions/sera-*.js, netlify/lib/*.js, src/hooks/useSeraInterview.js, src/pages/MeetSera.jsx and src/components/sera/*.
5. Never print secret values. Env var names are fine; values are not.
6. Anything listed under "Open decisions" (section 14) is the owner's call. Ask; do not assume.

1. The project in one page

YZI Works (Young Zone India Works) is a live platform connecting companies and candidates and preparing candidates to get placed. Its public site already contains Meet AI Sera, a free 5-minute AI voice interview that ends with an on-screen report and career roadmap. The owner now wants a second, much deeper product for institutes: Interview Arena.


- **What it is.** A realistic, multi-round interview simulator. Sera screens the student, then an HR BP, a functional interviewer and a final-round interviewer take over, all adapted to a company the student picked. Each round ends with feedback, levels unlock as the student improves, and reports are scoped to the student's current level.
- **Who uses it.** Students of partner institutes (candidate login) and institute staff (institute login). The institute dashboard shows each student's performance, status, scores, strengths and weak areas, and lets staff export reports.
- **Why institutes care.** A company loses money on a bad hire. An institute wins on placement numbers and reputation. So the pitch to institutes is better placement statistics and a readiness score they can show, not lower hiring risk.
- **Why it is not "just a prompt".** Any student can already ask a general chatbot to run a mock interview, so a good voice or an avatar is not a moat. Arena has to sell what a chat window cannot give: standardised scores institutes can compare and trust, progress over time, interviews built from real openings on YZI Works, a route to actual jobs, and institute-level data.
- **Outcome, not just simulation.** The owner does not want the product to stop at scoring. When a student falls short, it should name the specific gap (for example weak communication, or a skills gap against the

target CTC) and point to concrete learning, so the student feels cared for. Free diagnosis and a free first fix are non-negotiable; revenue comes from optional depth (section 10).

2. The two surfaces

Surface	Who	Where	Status
Meet AI Sera (open trial)	Anyone	Landing page, /meet-sera, existing repo	Live. Keep unchanged.
Interview Arena: candidate	Students of partner institutes	arena. subdomain (new)	To design, then build
Interview Arena: institute	Institute staff / placement cell	arena. subdomain (new)	To design, then build

A routing layer decides where a visitor goes: people arriving from an institute are sent to the arena subdomain; everyone else keeps using the landing-page trial. Eventually that layer needs a small change on the landing page. That is a change to production code, so it needs the owner's explicit approval when the time comes.



3. What exists today (read-only review of the repo)

3.1 Stack

React 19, Vite and Tailwind 4 on Netlify. Backend is Netlify Functions with Netlify Blobs as a simple key-value store (no traditional database). Cloudflare R2 for files, Retell for the voice call, OpenAI Responses API for analysis, Resend for email, Google Identity Services for Sera sign-in, MSG91 OTP for the older application forms. Repo conventions worth knowing: files use CRLF line endings, colour tokens are theme-based (dark default, light supported), and npm run build and npm run lint must pass before any commit.

3.2 How Meet AI Sera works

1. Google sign-in, then the candidate uploads one PDF résumé straight to R2 (presigned upload, 10 MB cap, PDF only).
2. sera-extract-resume sends the PDF to OpenAI in one call that confirms it is a résumé and returns first name, a highlight sentence, the field and the résumé text.
3. sera-start-call checks one-interview-per-email in Blobs, reserves the slot, then creates a Retell web call, passing candidate_name, resume_highlight and field as dynamic variables plus email and objectKey as metadata.
4. The browser runs the call through the Retell web SDK with a 5:00 session clock, a 35-second turn timer and stall detection. The UI is deliberately quiet: an orb, the clock and whose turn it is.
5. When the call ends, Retell calls sera-retell-webhook. It verifies the HMAC signature (X-Retell-Signature, five-minute freshness), and only for genuinely finished calls (agent_hangup or max_duration_reached) makes one OpenAI call to build the report.
6. The report holds strengths, growth areas, a Sera note, a three-tier roadmap (now, next, later) and video resources. Resource topics are resolved to real videos through the YouTube Data API. The webhook emails the team the report, transcript and a 7-day résumé link, and stores the report in Blobs.
7. The browser polls sera-get-report every two seconds (no LLM cost) and renders the stored report.

3.3 Locked mechanics of the trial (for context)

The owner's team fixed the trial as: greeting, one warm-up question, three easy skill questions (the third branches on the second), one combined close question (target role, salary band, timeline), hard stop at 5:00, two difficulty levels only. Sera never scores or explains the roadmap during the call. These rules belong to the trial. Arena's rounds are a new design and may differ, but the quiet UI and the cost discipline are worth inheriting.

3.4 Invariants worth carrying into Arena

- Analysis runs once per finished interview, inside the webhook, never from the browser.
- Calls that did not really finish skip the LLM entirely and are recorded honestly as incomplete.
- Reservations are released when a call fails to start.
- No unauthenticated endpoint may trigger paid API calls. An earlier analysis endpoint was deleted for exactly this reason.
- A turn timer limits answer length so a single answer cannot burn paid minutes.

3.5 Reusable patterns

File	What to reuse
netlify/functions/sera-start-call.js	Creating a Retell web call with dynamic variables and metadata. Arena: pass session and round ids plus a compact context summary.
netlify/functions/sera-retell-webhook.js	Signature check, disconnect-reason gate, single report call, resource lookup, email. Arena: same pattern per round, storing numeric scores.
netlify/lib/openai.js	generateJson({prompt, schema, file}) using the Responses API with strict JSON schema. Reuse for scoring, company profiles and the career-path step. The model name is hard-coded here; confirm it before scaling.
netlify/lib/seraReport.js	Report schema and prompt. Extend with per-round scores and transcript evidence.

File	What to reuse
sera-extract-resume.js	Résumé validation and highlight extraction in one call, passing the PDF to OpenAI as a file input. The file's comments explain why local PDF parsers crashed in the Netlify runtime; do not reintroduce them.
sera-create-upload.js, src/utils/seraUpload.js	Presigned direct upload to R2.
src/hooks/useSeraInterview.js	Retell client events, session and turn timers, stall detection, report polling.
src/components/sera/SeraOrb.jsx, orbMath.js	Particle orb ported from the MIT-licensed VoiceOrbs Particles Orb.
netlify/lib/eligibilityStore.js	Blobs store with an explicit siteID and token fallback, added because automatic injection did not work. Arena should use a real database instead.


3.6 Known gaps and risks found in the review

- sera-get-report returns the full report for any email posted to it, and sera-start-call trusts the email the browser sends without verifying the Google token server-side. Acceptable for a free trial; not acceptable for institute student data. Do not copy this pattern into Arena.
- Everything is keyed by email in Blobs: no history, no tenants. That cannot support levels, gating or cohort views.
- The trial report has no numeric scores, so gating and cohort analytics have nothing to work with.
- The report screen contains copy saying nothing is stored, while the report is kept in Blobs and the transcript and résumé link are emailed to admin. Check the wording (around line 218 of SeraReport.jsx). Consent and privacy copy matters much more once institutes are involved.
- YOUTUBE_API_KEY is not in the list of confirmed env vars in CLAUDE.md. If it is missing, video resources are silently dropped.
- sera-extract-resume and sera-create-upload are not gated by sign-in (documented, low cost).
- Deploy caveats from HANDOFF.md: a push does not automatically update the custom domain (someone else publishes), R2 CORS must allow the custom domain, and the Retell webhook URL must not be tied only to the netlify.app address.


3.7 Not verified (ask the owner)

- Whether the Retell prompt actually uses {{resume_highlight}}, and whether Retell's own maximum call duration is set to 300 seconds.
- Whether SERA_ADMIN_EMAILS is set in Netlify, and whether a monthly spend cap is set in the OpenAI dashboard.
- Status of report generation. CLAUDE.md (27 Aug) says the live call had never run end to end; the owner later reported that a live call connected and completed but the report step failed. The repo does not confirm that this was resolved.

4. Interview Arena: candidate experience



Pre-interview flow: after screening, career-path guidance appears, then the student picks a sector, sees three generated company profiles and chooses one, which locks the interview context.



Interview flow: three rounds in sequence, then a level-scoped report with resources.

4.1 Stages

Stage	Persona / voice	Purpose	Output
1. Screening	Sera (existing persona)	Baseline: communication, basic domain awareness, confidence	Screening result, feeds career path
2. Career path	Screen, no voice	2 or 3 realistic paths with reasoning and one video or course each. Shown only after screening so it rests on observed signal.	Student confirms a target role
3. Sector and company	Screen, no voice	Sector sub-list, then Sera researches and offers 3 company profiles; student picks one	Company context locked
4. HR BP round	HR persona, its own voice	Culture fit, background, salary expectations, communication	Score, feedback, notes for next persona
5. Functional round	Hiring-manager persona, its own voice	Role-specific and scenario questions, technical depth	Score, feedback, notes
6. Final round	Senior-leader persona, its own voice (optional video "boss round")	Pressure, judgement, negotiation, why hire you	Score, feedback, overall result
7. Report	Screen	Level-scoped report, gap diagnosis, resources, next-level check	Downloadable PDF

4.2 Sector sub-list (selection screen)

Sector	Sub-options
MNC / IT services	Product-based MNC, IT services and consulting, GCC (global capability centre)
Startup	Early stage (seed to Series A), growth stage (Series B and later), D2C / consumer
PSU / Government	Central PSU, state PSU, government-adjacent (banking, railways)
Custom	Student types a target company or industry

4.3 Company profile generation

- A separate backend step, not folded into the interview call. Input: sector, sub-option, role.
- Search the web for real companies matching the sector, gather basic facts (size, recent news, typical interview style if available), then have the LLM draft exactly three profiles, each distinct in size and tone.
- Cache results per sector, sub-option and role for a few days and refresh weekly, so each student does not trigger a fresh search.
- Label profiles as AI-generated from public sources and keep the source links. Wrong company facts would damage credibility with institutes.
- Build order: start with three hand-written profiles per sector, and add live research only after the whole loop works.

4.4 What makes a round feel real

- A stable persona per round: name, tone and voice, each with its own Retell agent.
- Company context in every prompt: name, size, tone, recent signals, interview style.
- Continuity: each round writes short notes for the next persona (for example "weak on Node.js in the HR round, probe this"), so the final interviewer can refer to earlier answers.

- Human pacing borrowed from the trial: react before moving on, one question at a time, at most one follow-up, graceful wrap-up language, gentle silence handling, and never cutting the student off mid-question.
- No scoring on the call. Scores and feedback come after the round.

4.5 Feedback, gating and levels

- After every round the student sees the key elements where they fell short, plus one strength, each tied to a specific moment in the answer. Vague labels such as "improve communication" are not enough.
- A round only unlocks the next one if the score clears that level's threshold. A failed round leads to the gap diagnosis, resources and a retry of that round, not forced progress.

Level	Difficulty	Unlock rule (placeholders to calibrate)
Beginner	Softer follow-ups, lower pass bar	Start here
Intermediate	Scenario questions, moderate bar	Average score across Beginner sessions above the threshold
Advanced	Hard scenarios, pressure, strict bar	Same rule on Intermediate

- Company type (MNC, startup, PSU) is independent of level. Full history is always stored; the student's report view shows only their current level.

4.6 Gaps and resources

Gap type	Signal	Fix
Communication	Pauses, filler words, clarity, answer structure	Free video, then practice drills and a retry of the round
Skill	CV and answers fall short of role expectations	Free playlist first, optional paid course after
Expectation	Salary expectation far from what the profile supports	Honest benchmarking content, not more interview practice

Resource lookup is its own module (gap type to resource library), kept out of the interview prompt so content can be curated without touching interview logic.

4.7 Resilience for home internet

- Voice is the default. It needs far less bandwidth than live video.
- A pre-check screen tests microphone and connection before any round starts.
- If the connection is weak, fall back automatically (video to voice) instead of failing.
- Save the context after every turn so a dropped student can resume the same round.
- Flag turns affected by a disconnect so they are not scored as hesitation, and use a more tolerant end-of-turn wait on slow connections.

5. Interview Arena: institute experience

Institute staff sign in with a separate institute login. They see only their own students.

Area	What it provides
Overview	Students invited, started and completed; average readiness; level distribution; improvement over time
Students	Filterable list by batch, level and status; last score; weakness tags; bulk invite

Area	What it provides
Student detail	Per-round scores, strengths, weak areas, trend across attempts, downloadable report
Batches	Create batches, upload a roster (CSV), generate invite links or codes
Analytics	Common weak areas across a cohort, before and after improvement, popular sectors and companies
Exports	CSV and PDF exports of student reports and batch summaries
Team and settings	Institute users and roles, privacy and consent settings

- Access is enforced on the server, not just hidden in the interface. Every record carries an `institute_id`.
- Log exports and views of student reports (audit trail).
- Decide what institutes may see: scores and summaries by default, transcripts or audio only with the student's explicit consent (section 14).

6. Data model (multi-tenant)

The trial's Blobs store cannot support this. Recommendation: a managed relational database (section 8).

Suggested tables:

Table	Key fields
institutes	id, name, slug, status, plan, created_at
institute_users	id, institute_id, email, role (admin or viewer), status
batches	id, institute_id, name, program, year
students	id, institute_id, batch_id, email, name, status (invited or active), consent_at, guardian_consent_ref
student_profiles	student_id, education, skills, projects, experience, cv_json, updated_at
interview_sessions	id, student_id, institute_id, level, sector, role, company_profile_json, context_json, status, started_at, ended_at
session_rounds	id, session_id, stage, persona, agent_id, retell_call_id, transcript_ref, score, rubric_json, strengths, gaps, status, duration
reports	id, session_id, level_at_generation, report_json, pdf_key, created_at
student_levels	student_id, level, progress_pct, sessions_at_level, avg_score, unlocked_at
gap_diagnoses	id, session_id, gap_type, evidence, resource_ids
resource_library	id, gap_type, level, title, url, type, is_paid, is_affiliate
company_profile_cache	id, sector, sub_option, role, profiles_json, generated_at, expires_at
consents	student_id, purpose, version, at, guardian_ref
audit_log	id, actor, action, target, at

7. The context object

One JSON object is built when the student picks a company and is read and updated by every round. It is what lets the final interviewer refer to what the HR interviewer heard.

```
{
  "session_id": "uuid", "institute_id": "uuid", "student_id": "uuid",
  "student_cv": { ...built from the student profile, no upload needed },
  "track": { "role": "Software Engineer", "sector": "MNC / IT services" },
  "company_profile": {
    "name": "...", "size": "large | mid | startup",
    "tone": "formal | fast-paced | protocol-heavy",
    "recent_signals": ["from web research"],
    "interview_style_notes": "...", "sources": ["urls"]
  },
  "level": "beginner | intermediate | advanced",
  "rounds": {
    "screening": { "status": "passed", "score": 71, "transcript_ref": "...", "flags": [] },
    "hr_bp":      { "status": "in_progress", "score": null, "transcript_ref": null, "flags": [] },
    "functional": { "status": "locked", "score": null, "transcript_ref": null, "flags": [] },
    "final":      { "status": "locked", "score": null, "transcript_ref": null, "flags": [] }
  },
  "running_notes_for_next_persona": "..."
}
```

8. Technology and APIs

Layer	Choice	Notes
Front end	React, Vite, Tailwind	Same family as the repo. Interface is prototyped first in Lovable (section 12); check what it exports before porting.
Voice	Retell, one agent per persona	Distinct voice per interviewer. Pass context through dynamic variables. Put the round's maximum duration in both the prompt and Retell's own setting.
LLM	OpenAI Responses API, strict JSON schema	Reuse generateJson. Keep the provider behind that function so it can change.
Company research	Search step, then LLM drafting	Search provider not chosen yet (section 14).
Database and auth	Managed Postgres with row-level security on institute_id	Recommendation only; owner to decide. Lovable integrates well with Supabase, but check data residency and DPDP implications.
Files	Cloudflare R2	Résumés, exported PDFs, recordings if kept.
PDF export	Server-side HTML to PDF	Store in R2, link from reports.
Email	Resend (existing)	Invites and notifications.
Hosting	Netlify, separate site for the subdomain	Keep independent of the live landing site.

8.1 Orb

The owner wants a more professional, standard orb and plans to use MIT-licensed open-source orbs from GitHub. The current site already ships a port of the VoiceOrbs Particles Orb (MIT). For any orb adopted: read its LICENSE file, keep the copyright notice, and test it on low-end phones, because particle systems can be heavy.

8.2 Video avatars

Real-time avatar services (Tavus, Anam, Beyond Presence, LemonSlice, LiveAvatar, HeyGen Interactive Avatar, D-ID Agents) are available to any competitor, so an avatar is a garnish, not a moat. Video needs much more

bandwidth than voice and adds per-minute cost. Recommendation: voice-first, with an optional video "boss round" for the final round only, falling back to voice on a weak connection.

9. Differentiators backlog

#	Feature	Why it is hard to copy
1	Readiness certificate institutes can share	Only meaningful with standardised scoring from one platform
2	Speech analytics from audio	Measured pace, filler words, pauses and answer length; a text prompt cannot measure these
3	Panel round	Two or three interviewers with different voices who interrupt and challenge
4	Indian placement context	Hinglish, campus-style questions, salary norms, PSU and government formats
5	Camera-based non-verbal feedback	Eye contact, posture, fidgeting; students will pay for this
6	Human mentor review of flagged sessions	A service institutes can be charged for

10. Monetisation

- Free for the student: the diagnosis and at least one genuine free resource per gap. Never put the diagnosis or the first fix behind payment.
- Institute side (preferred): a per-student per-session fee or a per-batch or semester licence, with the resource library bundled in. Students never see a paywall.
- Optional depth for direct users: affiliate links to paid courses for skill gaps, and the owner's own short paid modules for communication gaps.
- Rule: revenue logic must never influence which gap is flagged or which resource is recommended first.

11. Cost estimates (ballpark, unvalidated)

Item	Estimate	Notes
Voice, all-in (Retell incl. LLM, TTS, telephony)	About \$0.13 to \$0.20 per minute	Researched during planning; re-verify before quoting institutes
One full 4-stage session, about 25 to 30 minutes	About Rs 300 to 500	Assumes roughly Rs 85 per USD. Plan on about Rs 400. Higher than the earlier 3-round estimate because screening was added.
Report generation per session	Rs 4 to 12	Text only
500 sessions per month	About Rs 1.5 to 2.5 lakh	Before volume discounts
2,000 sessions per month	About Rs 6 to 10 lakh	Volume discounts should lower this
Video avatar minutes	Extra per-minute cost	Get a quote before pitching video

Build cost: earlier estimates (Rs 1.5 to 8 lakh, 11 to 13 weeks) assumed a submenu inside an existing dashboard. Two logins, an institute dashboard, a multi-tenant database and exports are a larger scope, so expect noticeably

more. Re-estimate after scoping against the repo. Institute pricing should sit comfortably above the running cost per session.

12. Recommended sequence (after the owner says go)

The owner will design the interface first with remaining Lovable credits, using mock data, and connect the backend afterwards. Sequence:

Phase	Scope
0. Interface (owner, Lovable)	All screens in section 13 with realistic mock data and every state (loading, error, incomplete, poor connection). Export to GitHub.
1. Foundation	New repo and subdomain, database schema, candidate and institute authentication (server-verified), institute onboarding, roster and invites.
2. Candidate core	Profile to CV JSON, screening agent and webhook, career path, sector pick, hand-written company profiles, context object.
3. Rounds	HR, functional and final agents with distinct voices, scoring rubric, per-round feedback, gating, levels.
4. Reports	Level-scoped report, PDF export, gap diagnosis and resource library.
5. Institute dashboard	Live data, filters, cohort analytics, exports, audit log.
6. Hardening	Live company research, resilience features, security review, privacy and consent, load test.
7. Pilot	One institute, agreed success metrics (readiness improvement, placement outcomes).

Start with one thin vertical slice: one institute, one student, one company profile, three rounds, one report, one institute view. Expect a rough working slice in days or a week, but not the whole product in a day. Timelines should be re-estimated by whoever scopes the repo.

13. Screen inventory (for Lovable and later mapping)

13.1 Entry and login

- Institute entry: arrive via institute link, or enter an institute code.
- Candidate login, institute login, accept-invite and password or link reset.

13.2 Candidate screens

#	Screen	Key elements
1	Home	Level ring and progress, next-step call to action, latest report
2	My profile	CV builder: education, skills, projects, experience; completeness meter
3	Screening	Microphone and connection pre-check, then live conversation with orb
4	Career path	2 or 3 suggested paths with reasoning and one resource each
5	Choose sector	Sector cards, then the sub-list
6	Choose company	Three AI-generated company cards with source links
7	Briefing	Interviewer introduction, what to expect, ready check
8	Live interview	Orb, clock, whose turn it is, connection status. Quiet interface, no live transcript.

#	Screen	Key elements
9	Round feedback	Key gaps and one strength with evidence, pass or retry, next step
10	Levels and progress	Beginner to Advanced ladder, unlock rules, locked states
11	Report	Level-scoped scorecard, gaps, roadmap, download PDF
12	Resources	Free and paid clearly marked, grouped by gap
13	History	Past attempts and trend
14	Privacy and consent	What the institute can see, data deletion request

13.3 Institute screens

#	Screen	Key elements
1	Overview	KPIs, level distribution, improvement chart, recent activity
2	Students	Table with filters, status, level, last score, weakness tags; bulk invite
3	Student detail	Round scores, strengths, weaknesses, trend, report download
4	Batches	Create batch, upload roster, invite links or codes
5	Analytics	Common weak areas, before and after, sector and company popularity
6	Exports	CSV and PDF exports for students and batches
7	Team and settings	Users, roles, privacy settings

Design direction: premium and calm; candidate side mobile-first (students are on phones), institute dashboard desktop-first; light and dark themes like the current site. The interactive demo from earlier planning (<https://claude.ai/artifact/Aa1EYstaVbDGw1jiU96rma>) is a visual reference only; it shows the older idea of a submenu inside a dashboard, and the structure has since changed to the standalone arena.

14. Open decisions (owner's call)

#	Decision	Recommendation / note
1	How institute students reach the arena	Institute-specific link or code plus an uploaded roster. Do not trust a plain query parameter.
2	Who creates institute accounts	YZI admin only at first; self-serve later.
3	Authentication	Candidates: Google or email OTP with server-side token verification. Institute users: email with magic link or password.
4	Database and backend	Managed Postgres with row-level security; confirm residency and DPDP fit.
5	Where the student profile and CV come from	The owner says the YZI Works dashboard holds profiles and that he built it; HANDOFF.md says another person owns "the dashboard". Clarify which codebase and API before any integration. Until then, use an arena-native profile.
6	What institutes can see	Scores and summaries by default; transcripts or audio only with explicit student consent.

#	Decision	Recommendation / note
7	Minors and consent	Some students may be under 18. India's DPDP Act requires verifiable parental consent for children's data. Get legal advice before onboarding such batches.
8	Data retention and recordings	Set retention and deletion windows; decide whether audio is kept at all.
9	Subdomain name and institute branding	Name to be chosen; white-label logo and colours later.
10	Avatar or video	Voice-first; optional final-round video later.
11	Languages	English and Hinglish first; regional languages later.
12	Company research provider	Search provider and a review policy for AI-generated company profiles.
13	Scoring rubric ownership	Owner plus an HR domain expert define it; calibrate on real sessions and audit periodically.
14	Pricing model	Per student per session versus per batch or semester licence.
15	First pilot institute and success metrics	Agree readiness improvement and placement targets up front.

15. Risks to keep in view

- Scope versus timeline: the owner hopes to move very fast, but the full product is much larger than the trial. Insist on the thin vertical slice.
- Scores institutes will rely on must be defensible. An uncalibrated AI score presented as objective truth is a reputational risk; sample real sessions against human judgement.
- Student data: multi-tenant isolation, consent, minors and retention all need designing in from the start, not adding later.
- Do not copy the trial's weak authentication and email-keyed storage into Arena.
- AI-generated company facts can be wrong; show sources and cache reviewed results.
- Vendor pricing and model availability change. Re-check Retell, OpenAI and any avatar service before quoting institutes.
- Video adds cost and connection problems for students on home internet.

16. Your first steps

8. Read the repo in the order given in section 0. Do not change anything.
9. Check the items in section 3.7 with the owner.
10. Summarise back to the owner, in your own words, what Arena is and how it differs from the trial, and say what you are unsure about.
11. Ask the questions in section 14, most important first (1, 3, 4, 5, 7).
12. Propose a scoped plan and estimate for the thin vertical slice. Wait for an explicit go before writing any code.

17. Earlier deliverables

Earlier planning produced a first PDF (architecture and costing for a submenu inside the dashboard), a v1 technical spec PDF, and the interactive submenu demo linked above. This document supersedes them wherever they disagree, mainly because the product is now a standalone subdomain with two logins, and the scope, data model and cost picture have grown accordingly.